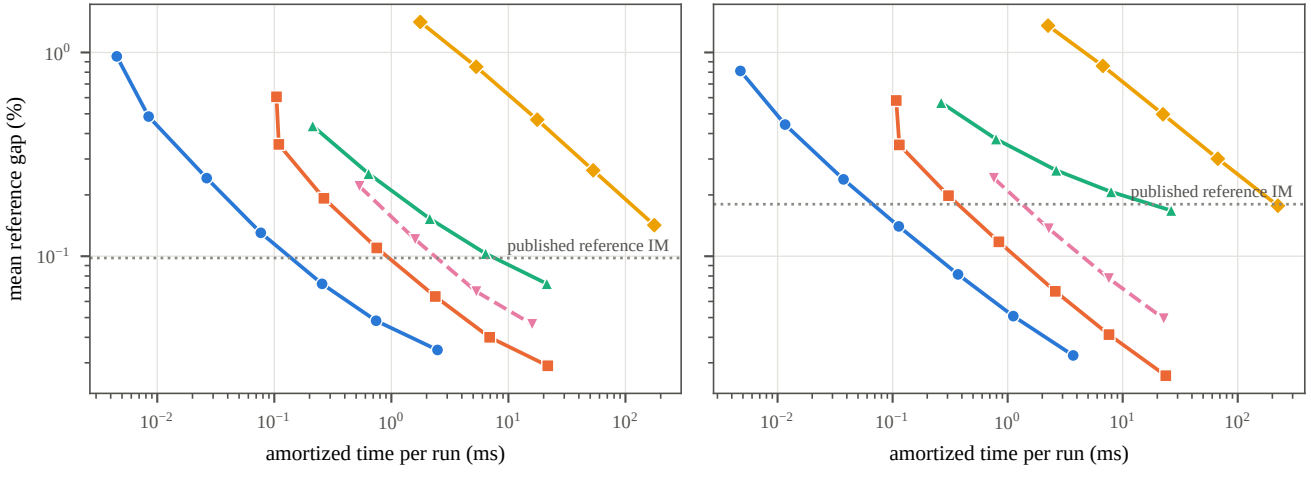
(a) Max-3-Cut, mean of G1–G5
(b) Max-4-Cut, mean of G1–G5
mean reference gap (%)
10^{0}
10^{-1}
published reference IM
10^{-2}
10^{-1}
10^{0}
10^{1}
10^{2}
amortized time per run (ms)
PdSB (fused CUDA)
PdSB (PyTorch)
Potts SA (136 CPU threads)
one-hot dSB
reference one-hot IM, our run (G1 only)